

Algorithmic Game Theory

Notes by Arjun Maneesh Agarwal

based on course by Prajakta Nimbhorkar

1. Games and Equilibria : We are assuming games where everyone is acting selfishly.
When everyone doesn't wish to switch strategies, we get an equilibrium.
 - We will see two-player games and multiplayer games.
 - Equilibrium and existence and computation.
2. Mechanism Design : Given multiple selfish agents, we need to make a mechanism that encourages (or discourages using hardness) the behavior we wish to encourage.
 - We will see auctions, voting etc.
3. Fair Division: This is an allocation problem where we want to allocate some resources in a 'fair' (or less unfair) way.
 - Distributing divisible and indivisible resources fairly among agents, given valuation.
 - EF, EQ, PROP and its relaxations.

Table of Contents

1. Games and Equilibria	3
1.1. Braess's Paradox	3
1.2. Prisoner's Dilemma	4
1.3. Pollution Control	6
1.4. RPS	6
1.5. Sealed Bids Auctions	7
1.6. Battle of the Sexes	8
1.7. Properties of MSNE	9
2. Equilibrium Computation	10
2.1. Computing a SDSE	10
2.2. Two Player Zero Sum Games	12
2.3. Linear Programming for the the win	14
2.3.1. Duel of LP	14
2.4. Strong Duality	15
2.5. Existence of Nash Equilibrium (mixed) in finite strategic form games	15
2.6. Two Player Non-Zero Sum Games	17
2.6.1. NE by Support enumeration	17
2.6.2. Lemke-Howson Algorithm	18
2.7. Mixed Nash Equilibrium in 2 Player Games	20
2.8. Complexity of algorithm	21
2.9. Alternatives to Nash Equilibrium	23
2.9.1. Corelated Equilibria	24
3. Mechanism Design	26
4. Fair Division	26
4.1. Detour: Sperner's Lemma	27
4.2. Back to Fair Division	30
4.2.1. Cake-Cutting Problem	30
4.2.2. Harmounious Rent Problem	32
4.3. Fair Division of indivisible goods	34


1. Games and Equilibria

We make the following assumptions:


1. Rationality/selfishness : Each agent attempts to maximize own payoff and believes others will do so as well as well has knowledge of same behavior being
2. Intelligence: Agents have enough computational resources to take into account the strategies and behaviors of others.
3. Common Knowledge: P, S_i, μ are known to everyone and everyone knows that everyone knows them and so on.

1.1. Braess's Paradox

Taking the example of traffic, let the network be:



In this case, the equilibrium is 1.5 hours as traffic is divided in the two paths. Let's say the authorities build a bridge:



This will lead to a (weak) equilibrium of 2 hours as all agents take the route s-v-u-t.

This is called Braess's Paradox.

Side Note

We have actually seen this example occur almost verbatim in Seoul where the demolition of a bridge in 2005, solved a 2 decade old traffic problem.

Furthermore, Steinberg and Zangwill, 1983; showed that in a random network, Braess's paradox is approximately 50% likely to occur in an random network.

It is also a consideration in google maps where the navigation sometimes suggests a longer route, not to avoid a pre-existing traffic jam, but to avoid creating one. (This is why google map is not referentially transparent and may suggest two people different routes between same places at same time)

We also would like to define the term 'cost of anarchy' here which refers to the ratio of equilibrium payouts (or cost) divided by optimal payouts(or cost). In this case, anarchy was $\frac{2}{1.5} = \frac{4}{3}$ times costlier.

1.2. Prisoner's Dilemma

		P2	
		<i>C</i>	<i>D</i>
P1	<i>C</i>	4, 4	5, 1
	<i>D</i>	1, 5	2, 2

Here both players have incentive to deviate from the mutually best case as confessing (strongly) dominates denial and hence, is a Nash equilibrium.

Let's define these words formally.

Definition: Game

To define a game we need to know the players and strategies for each player. Thus, a game $G = (P, S, \mu)$ where for i in P , S_i in S denotes the set of possible strategies for player i .

In a particular play of the game denoted as s , i plays strategy s_i forming the strategy profile or game vector (s_1, s_2, \dots, s_n) .

We denote the payoffs or costs of the game for player i using a function $\mu_i : S_1 \times S_2 \times \dots \times S_n \rightarrow \mathbb{R}$ which takes the game vector and tells the payoff for player i . $\mu : S_1 \times S_2 \times \dots \times S_n \rightarrow \mathbb{R}^n; \mu(s) = (\mu_1(s), \mu_2(s), \dots, \mu_n(s))$ is the payoff vector returning function.

Finally, for conciseness, we sometimes denote the choices for all players except i as s_{-i} . Thus, $\mu_i(s_i, s_{-i})$

Definition: Strongly Dominant Strategy

Player i 's strategy \hat{s}_i is a **Strongly Dominant Strategy** to the strategy s_{-i} of other players if:

$$\mu_i(\hat{s}_i, s_{-i}) > \mu_i(s'_i, s_{-i}) \forall s'_i \in S_i$$

Definition: Weakly Dominant Strategy

Player i 's strategy \hat{s}_i is a **Weakly Dominant Strategy** to the strategy s_{-i} of other players if:

$$\mu_i(\hat{s}_i, s_{-i}) \geq \mu_i(s'_i, s_{-i}) \forall s'_i \in S_i$$

and

$$\exists s'_i \in S_i \text{ s.t. } \mu_i(\hat{s}_i, s_{-i}) > \mu_i(s'_i, s_{-i})$$

Definition: Weakly Dominant Strategy / Best Response(BR)

Player i 's strategy \hat{s}_i is a **best response** to the strategy s_{-i} of other players if:

$$\mu_i(\hat{s}_i, s_{-i}) \geq \mu_i(s'_i, s_{-i}) \forall s'_i \in S_i$$

or if \hat{s}_i solves $\max_{S_i} \mu_i(s_i, s_{-i})$.¹

Definition: Strongly Dominated Strategy Equilibrium

Informally, this is when all players are playing dominant strategies.

It is a strategy profile $(s_1^*, s_2^*, \dots, s_N^*)$ is a SDSE if for each i , her choice s_i^* is strongly dominant to all .

Definition: Nash Equilibrium(NE)

Informally, this is when all players are playing BR to each other.

Formally, it is a strategy profile $(s_1^*, s_2^*, \dots, s_N^*)$ is a NE if for each i , her choice s_i^* is a best response to s_{-i}^* .

¹Sometimes we can replace the actual strategies by others by p , which is i 's belief of others' choices. This is often done in econ, but here, one of our assumptions (Common knowledge) allows us to not need it. The BR word is also from Econ.

Remark

The reason Nash equilibrium is important is as it leads to no-regrets as no individual can do strictly better by deviating holding others fixed. Second, as it is a self-fulfilling prophecy (we will see what this means).

1.3. Pollution Control

Let there be n countries with the cost of pollution control being 3. If a country doesn't control pollution, it adds cost 1 to each country.

This leads to the dark case of no one controlling pollution. We can check that if only k countries control pollution, all have incentive to not control.

Definition

A dominant strategy is what is best for a player regardless of whatever strategy the rest of the players choose to adopt.

1.4. RPS

		P2		
		<i>R</i>	<i>P</i>	<i>S</i>
P1	<i>R</i>	0, 0	1, -1	-1, 1
	<i>P</i>	1, -1	0, 0	-1, 1
	<i>S</i>	0, 0	-1, 1	1, -1

There is no pure Nash equilibrium here. But allowing randomized strategy, we can get an equilibrium.

Definition: Mixed Strategy

A mixed strategy p_i is a randomization over i 's pure strategies.

- $p_i(s_i)$ is the probability i assigns to the pure strategy s_i .
- $p_i(s_i)$ could be zero for some i .
- p_i could be one (in case of a pure strategy).

Definition: Mixed Strategy Profile

A mixed strategy profile is a tuple (p_1, p_2, \dots, p_n) where each p_i is a mixed strategy for player i . It represents the joint behavior of all players where each one plays a randomized strategy independently.

Definition: Expected Payoff

The expected payoff for player i is:

$$\sum_{(s_1, \dots, s_n) \in S_1 \times \dots \times S_n} \mu_i(s_i, s_{-i}) \prod_{a=1}^n p_a(s_a)$$

Definition: Mixed Strategy Nash Equilibrium

A mixed strategy profile $(p_1^*, p_2^*, \dots, p_n^*)$ is a mixed strategy NE (MSNE) if for each player i , p_i^* is a BR to p_{-i}^* . That is:

$$\mu_i(p_i^*, p_{-i}^*) \geq \mu_i(\sigma_i, p_{-i}^*) \forall \sigma_i \in \text{probability distribution over } S_i$$

Nash's Theorem 1.1. *Every bimatrix game has a (pure or mixed) Nash equilibrium.*

Theorem 1.2. *For zero-sum games, it is polytime; otherwise, it is PPAD-hard.*

PPAD-hard $\stackrel{?}{\Rightarrow}$ NP-Hard \iff NP $\stackrel{?}{=}$ co-NP. This is open and unknown.

1.5. Sealed Bids Auctions

There is one seller and n buyers and one item that the seller wants to sell.

Buyers have non-negative valuations over this item v_i .

The rules of this auction are that each buyer submits a sealed bid b_i be the bid amounts. The highest bid gets the item (in case of ties, by the lower index).

If buyer i gets the item, and pays t_i , then their $\mu_i = v_i - t_i$ and payoff for other buyers is 0.

Definition: First Price Auctions

In first price auction, $t_i = b_i$.

Definition: Second Price Auction

Buyer who gets the items (say i) pays the second highest bid

Theorem 1.3. *Second price auction is weakly revealing (that is it is a weakly dominant strategy to bid truthfully).*

Proof. In such an auction setting (regret less, envy less, fee less auctions), our payoff is zero unless we win.

What agents are hence trying to avoid is the so called ‘winner’s curse. Hence, we can restrict the analysis to bidding as if we know we will win.

In that case, bidding your value is weakly dominant as if everyone is bidding their value; you can’t do better by

- bidding lower than your value, $\tilde{v} < v$, as another agent may have valuation v^* such that $\tilde{v} < v^* < v$ and will now win instead of you.
- bidding above your valuation, $\tilde{v} > v$, as another agent may have valuation v^* such that $\tilde{v} > v^* > v$ and you will win instead of them; However, with the winners curse.

■

Side Note

This is the only weakly dominant equilibrium. Although, this is by far not the only nash equilibrium. $(v_1, 0, \dots, 0)$ is a NE; $(v_2, v_1, 0, \dots, 0)$ is a NE; and so on.

This is called a Bayesian Nash Equilibrium as given our beliefs of the nature of the world, this maximizes our payoff. (Here the belief is that all valuations (in some radius around ours) are likely and that we may win).

1.6. Battle of the Sexes

Definition

We represent a situation where two agents must simultaneously take an action where each of them prefers one option over the other but prefer coordination over doing different things. And example occurrence would be²:

		Wife	
		Cricket	Music
Husband	Cricket	2,1	0,0
	Music	0,0	1,2

We can try to brute force this. Let’s look at the payoffs for both the players

$$\begin{aligned}
 \mu_1(\sigma_1, \sigma_2) &= \sigma_1(C)\sigma_2(C)\mu_i(C, C) + \\
 &\quad \sigma_1(C)\sigma_2(D)\mu_i(C, D) + \\
 &\quad \sigma_1(D)\sigma_2(C)\mu_i(D, C) + \\
 &\quad \sigma_1(D)\sigma_2(D)\mu_i(D, D) \\
 &= 2\sigma_1(C)\sigma_2(C) + \sigma_1(D)\sigma_2(D) \\
 &= 2\sigma_1(C)\sigma_2(C) + (1 - \sigma_1(C))(1 - \sigma_2(C)) \\
 &= 3\sigma_1(C)\sigma_2(C) - \sigma_1(C) - \sigma_2(C) + 1
 \end{aligned}$$

²Which is extremely stereotypical and doesn’t represent the author and hopefully the prof’s views.

And similarly

$$\mu_2 = 3\sigma_1(C)\sigma_2(C) - 2\sigma_1(C) - 2\sigma_2(C) + 2$$

Using $\mu_1(\sigma_1^*, \sigma_2^*) \geq \mu_1(\sigma_1, \sigma_2^*) \forall \sigma_1 \in \Delta(S_1)$

$$\sigma_1^*(C)[3\sigma_2(C) - 1] \geq \sigma_1(1)[3\sigma_2(C) - 1] \forall \sigma_1(c) \in (0, 1)$$

Similarly

$$\sigma_2^*(C)[3\sigma_1(C) - 2] \geq \sigma_2(1)[3\sigma_1(C) - 2] \forall \sigma_2(c) \in (0, 1)$$

If $3\sigma_2^*(C) - 1 > 0 \Rightarrow \langle (1, 0), (0, 1) \rangle$.

If $3\sigma_2^*(C) - 1 < 0 \Rightarrow \langle (0, 1), (1, 0) \rangle$.

If $3\sigma_2^*(C) - 1 = 0 \Rightarrow \langle (\frac{2}{3}, \frac{1}{3}), (\frac{1}{3}, \frac{2}{3}) \rangle$.

1.7. Properties of MSNE

Expected payoff for player i will be:

$$\begin{aligned} \mu_i(p_1, p_2, \dots, p_n) &= \sum_{(s_1, \dots, s_n) \in S_1 \times \dots \times S_n} \mu_i(s_i, s_{-i}) \prod_{a=1}^n p_a(s_a) \\ &= \sum_{s_i \in S_i} \sum_{s_{-i} \in S_{-i}} p_i(s_i) p_{-i}(s_{-i}) \mu_i(s_i, s_{-i}) \\ &= \sum_{s_i \in S_i} p_i(s_i) \sum_{s_{-i} \in S_{-i}} p_{-i}(s_{-i}) \mu_i(s_i, s_{-i}) \\ &= \sum_{s_i \in S_i} p_i(s_i) \mu_i(s_i, p_{-i}) \end{aligned}$$

This makes the utility of a player a convex combination of their pure strategy payoffs.

Definition: Convex Combination

A convex combination of a_1, a_2, \dots, a_n is $\sum_{i=1}^n \lambda_i a_i$ where $\lambda_i \in [0, 1]$ and $\sum_{i=1}^n \lambda_i = 1$.

An obvious observation is that convex combination of $a_1, \dots, a_n \leq \max\{a_i\}$.

Which implies that the payoff with mixed strategy is less than equal to max payoff with a pure strategy s_i . This implies:

$$\max_{\sigma_i \in \Delta(S_i)} \mu_i(\sigma_i, \sigma_{-i}) = \max_{s_i \in S_i} \mu_i(s_i, \sigma_{-i})$$

Theorem 1.4. $(\sigma_1^*, \dots, \sigma_n^*)$ is a MSNE if and only if $\forall i \in N$

1. $\mu_i(s_i, \sigma_{-i}^*)$ is the same for all $s_i \in \text{support of } \sigma_i^*$.
2. $\mu_i(s_i, \sigma_{-i}^*) \geq \mu_i(s_i', \sigma_{-i}^*) \forall s_i' \notin \text{support of } \sigma_i^*, s_i \in \text{support of } \sigma_i^*$.

Remark: Implications

Each player gets the same payoff for any pure strategy in the support of the MSNE strategy.

2. Equilibrium Computation**2.1. Computing a SDSE****Example**

		P_2		
		D	E	F
P_1	A	4,3	5,1	6,2
	B	2,1	8,4	3,6
	C	3,0	9,6	2,8

Notice, E is dominated by F as $\forall i, \mu_2(F, s_{-i}) > \mu_2(E, s_{-i})$. This practically makes the game:

		P_2	
		D	F
P_1	A	4,3	6,2
	B	2,1	3,6
	C	3,0	2,8

Now A dominates B, C and after elimination, D will dominate F and we will be left with the nash equilibrium of 4, 3.

Also note, (A, D) is a PSNE.

Note, this algorithm by no means gives all the nash equilibrium, if there are multiple. We don't run into such issues with SDSE as only one can exist by the definition of Strict dominance.

Example

		P_2		
		D	E	F
P_1	A	3,1	0,1	0,0
	B	0,1	4,1	0,0
	C	1,1	1,1	5,0

We can eliminate F as D, E are strictly better.

		P_2	
		D	E
P_1	A	3,1	0,1
	B	0,1	4,1
	C	1,1	1,1

Here, we can eliminate C as if P_2 plays D or E , we are better off playing A or B respectively.

Another way to argue the same is $\mu_1\left(\frac{A+B}{2}, s_{-i}\right) > \mu_1(C, s_{-i})$.

Side Note

Here, playing C is called the Bayesian Equilibrium as it is safe.

Example

		P_2		
		X	Y	Z
P_1	A	7,7	4,2	1,8
	B	2,4	5,5	2,3
	C	8,1	3,2	0,0

We can get the gurentee finding the nash equilibrium by identifying the highest entry per row and column wrt the respective players.

		P_2		
		X	Y	Z
P_1	A	<u>7</u> ,7	4,2	1, <u>8</u>
	B	2,4	<u>5</u> , <u>5</u>	2,3
	C	<u>8</u> ,1	3, <u>2</u>	0,0

Making (5, 5) the PSNE.

Algorithm 2.1. For the column player, mark the maximum entry for them.

For the row player, mark the maximum entry for them.

The intersections are PSNEs.

2.2. Two Player Zero Sum Games

Definition: Two Player Zero Sum Games

A game with:

- $N = \{1, 2\}$
- $S = \{S_1, S_2\}$
- $\mu = \{\mu_1, \mu_2\}$
- $\mu_1(s_1, s_2) = -\mu_2(s_1, s_2)$

In these games, we only need to specify the payoffs for one player. We, by convention, do this for the row player.

B

A

2	1	1
-1	1	2
1	0	1

For any strategy, i of P_1 and P_2 will choose a strategy such that,

- P_1 chooses $\max_i \min_j a_{ij}$ (maxmin)
- P_2 will choose $\min_j \max_i a_{ij}$. (minmax)

If maxmin and minmax are equal, we are done and the value is the nash equilibrium.

Definition: Maxmin-Minmax

The maxmin value refers to $\max_{i \in S_1} \min_{j \in S_2} a_{ij}$.

The minmax value refers to $\min_{j \in S_2} \max_{i \in S_1} a_{ij}$.

If a PSNE exists, we will get it by this process.

If there is no PSNE, this obviously doesn't work.

$P2$

$P1$

	R	P	S
R	0	-1	1
P	1	0	-1
S	-1	1	0

We can also discuss Saddle points.

Definition: Saddle Point

(i, j) is a saddle point of a matrix A if $a_{ij} \geq a_{kj} \forall k$ and $a_{ij} \leq a_{il} \forall l$

Theorem 2.2. *If i, j and k, h are saddle points, then (i, h) and k, j are also saddle points.*

Proof.

$$a_{ij} \geq a_{kj} \geq a_{kh} \geq a_{ih} \geq a_{ij}$$

And we are done by squeeze theorem. ■

Definition: Mixed Strategy in 2 player zero sum game

Let $|S_1| = m$ and $|S_2| = m$. Let $x = (x_1, x_2, \dots, x_n)$ be a mixed strategy for P_1 and $y = (y_1, y_2, \dots, y_n)$ be a mixed strategy for P_2 .

The expected payoff is $\sum_{i=1}^n \sum_{j=1}^m x_i y_j a_{ij} = x^T A y$

We can define maxmin and minmax values here as:

Definition: Maxmin-minmax

maxmin value = $\max_{x \in \Delta(S_1)} \min_{y \in \Delta(S_2)} x^T A y$

minmax value = $\min_{y \in \Delta(S_2)} \max_{x \in \Delta(S_1)} x^T A y$

Lemma 2.3.

$$\min_{y \in \Delta(S_2)} x^T A y = \min_{j \in S_2} \sum_{i=1}^m a_{ij} x_i$$

Proof. It is obvious that

$$\min_{y \in \Delta(S_2)} x^T A y \leq \min_{j \in S_2} \sum_{i=1}^m a_{ij} x_i$$

as $S_2 \subseteq \Delta(S_2)$.

To do the other way round,

$$\begin{aligned} x^T A y &= \sum_{j=1}^n y_j \sum_{i=1}^m x_i a_{ij} \\ &\geq \sum_{i=1}^n y_j \min_{k \in S_2} \sum_{i=1}^m x_i a_{ik} \\ &= \min_{k \in S_2} \sum_{i=1}^m x_i a_{ik} \cdot \sum_{i=1}^n y_j \\ &= \min_{k \in S_2} \sum_{i=1}^m x_i a_{ik} \end{aligned}$$

Thus, by squeeze theorem, we are done. ■

2.3. Linear Programming for the the win

We make an LP for the row player as:

The row player has to to maximize $\min_{j \in S_2} \sum_{i=1}^m a_{ij}x_i$ with the constrains $\sum_{i=1}^m x_i = 1, x_i \geq 0 \forall i \in \{1, 2, \dots, m\}$.

Equivalently, maximize z with the constrains $z \leq \sum_{i=1}^m a_{ij}x_i \forall j \in \{1, 2, \dots, n\}, \sum_{i=1}^m x_i = 1, x_i \geq 0 \forall i \in [m]$.

Exercise : Column Player's LP

Define the column players LP.

Solution. minimize w such that $w \geq \sum_{j=1}^n a_{ij}y_j \forall i \in [m], \sum_{j=1}^n y_j = 1, y_j \geq 0$ for all $j \in [n]$. ■

2.3.1. Duel of LP

Example : Toy LP

max $3x_1 + 2x_2 + x_3$ with constraints

1. $x_1 + x_2 + x_3 \leq 2$
2. $3x_1 + x_3 \leq 4$
3. $x_1 + x_2 + x_3 \leq 5$
4. $x_1, x_2, x_3 \geq 0$

We can get an upper bound by $3 * \text{eq 3} \Rightarrow 3x_1 + 3x_2 + 3x_3 \leq 15$.

We can get a better upperbound by $\text{eq 1} + \text{eq 2} \Rightarrow 3x_1 + 3x_2 + 3x_3 \leq 9$.

The idea is that as atleast one coefficient is bigger, the sum is bigger and is an upperbound.

But doing this for all combinations is going to be painfully time consuming. Let's try to do this algebraically and let the coefficients of equation 1,2,3 be y_1, y_2, y_3 in the linear combination.

$$x_1(y_1 + 3y_2 + y_3) + (y_1 + y_3)x_2 + (y_2 + y_3)x_3 \leq 2y_1 + 4y_2 + 5y_3$$

This leads to a LP

Minimize $2y_1 + 4y_2 + 5y_3$ with the constraints:

- $y_1 + 3y_2 + y_3 \geq 3$
- $y_1 + y_3 \geq 2$
- $y_2 + y_3 \geq 1$

Theorem 2.4. *The Row player LP and Column player LP are dual pairs.*

Proof. DIY. It's just transposing stuff here and there. ■

2.4. Strong Duality

Optimum value of f is the optimal value of its primal dual.

Ler $x^* = (x_1^*, x_2^*, x_3^*, \dots, x_m^*)$, z^* be the optimum values for Row player.

z^* must attain equality at some $j^* \in [n]$ that is $z^* = \sum_i a_{ij^*} x_i^*$ and $z^* \leq \sum_i a_{ij} x_i^* \forall j \in [n]$.

By lemma, $z^* = \min_{j \in S_2} \sum_{i=1}^n x_i^* = \min_{y \in \Delta(S_2)} x^{*T} A y$ and similarly, $w^* = \max \sum_{j=1}^n a_{ij^*} y_j = \max_{x \in \Delta(S_2)} x^T A y^*$.

Thus,

$$\min_{y \in \Delta(S_2)} x^{*T} A y = \max_{y \in \Delta(S_2)} x^T A y^*$$

This is a nash equilibrium as no player can unilaterally increase payoff by moving.

2.5. Existence of Nash Equilibrium (mixed) in finite strategic form games

Ler $(N, \langle S_i \rangle, \langle \mu_i \rangle)$ be the game where $N = \{1, 2, \dots, n\}$, S_i is strategy set for player i , with $|S_i| = m \forall i$ and $\mu_i : S_i \times S_{-i} \rightarrow \mathbb{R}$ is the payoff function for player i .

Mixed strategy σ_i is a probability distribution over S_i .

Δ_i of $\Delta(S_i)$ denotes the set of all mixed strategies for player i . $\Delta = \Delta_1 \times \Delta_2 \times \dots \times \Delta_m$.

Expected payoff from δ_i for player i is

$$\mu_i(\sigma_i) \triangleq \sum_{(s_1, \dots, s_n) \in S_1 \times S_2 \times \dots \times S_n} \left(\prod_{j=1}^n \sigma_j(s_j) \right) \mu_i(s_1, \dots, s_n)$$

Definition

σ^* is a Nash Equilibrium if for every player i ,

$$\mu_i(\sigma^*) \geq \mu_i(s_i, \sigma_{-i}^*) \forall s_i \in S_i$$

Brouwer's Fixed Point Theorem 2.5. Let B be a closed, bounded convex set. Let $f :: B \rightarrow B$ be a continuous function, then $\exists x \in B$ s.t. $f(x) = x$

We want to define B, f such that the Nash Equilibrium is the fixed point of f . Define $B = \Delta$.

Define f such that $\sigma \in \Delta$ is not a NE then $f(\sigma) \neq \sigma$ NS IF σ^* is a NE then $f(\sigma^*) = \sigma^*$.

Attempt 1: Define $f(\sigma) = \rho$ where $\rho \in \Delta$, $\rho = (\rho_1, \rho_2, \dots, \rho_n)$ such that

$$\rho_i = \arg \max_{\delta_i \in \Delta_i} \mu_i(\sigma_{-i}, \delta_i)$$

Recall, this is called the best response.

The problem is, this is not a function and if we apply a tie breaking rule, it is not continuous. We could use the powerset as the domain, but that would require Katakumi's fixed point which is more analytical than we wish to be.

Example : Matching Pennies

		P_2	
		H	T
P_1	H	1, -1	-1, 1
	T	-1, 1	1, -1

Let row player choose the mixed strategy $(p, 1-p)$ and column player choose $(q, 1-q)$. The problem is

$$p > \frac{1}{2} \Rightarrow q = 0$$

$$p < \frac{1}{2} \Rightarrow q = 1$$

$$p = \frac{1}{2} \Rightarrow 0 \leq q \leq 1$$

This is an counter example to the continuity of f .

Modifying f : We first define a gain function of player i for deviation to s_{ij} from σ_i .

$$g_{ij} = \max\{\mu_i(\sigma_{-i}, s_{ij}) - \mu_i(\sigma), 0\}$$

We now define

$$f_{ij}\sigma = \frac{\sigma_{ij} + g_{ij}(\sigma)}{\sum_{k=1}^m (\sigma_{ik} + g_{ik}(\sigma))} = \frac{\sigma_{ij} + g_{ij}(\sigma)}{1 + \sum_{k=1}^m g_{ik}(\sigma)}$$

We want

$$\sum_{k=1}^m f_{ik}(\sigma) = 1$$

We will define

$$\begin{aligned} f(\sigma) &:= \rho \\ f(\sigma) &= f(\sigma_{11}, \dots, \sigma_{1m}, \dots, \sigma_{n1}, \dots, \sigma_{nm}) \\ &= (\rho_1, \dots, \rho_{1m}, \dots, \rho_{n1}, \dots, \rho_{nm}). \end{aligned}$$

Let σ^* be a fixed point of f .

$$\begin{aligned}
f_{ij} &= \sigma_{ij}^* \\
\sigma_{ij}^* &= \frac{\sigma_{ij} + g_{ij}(\sigma)}{1 + \sum_{k=1}^m g_{ik}(\sigma)} \\
&\iff g_{ij}(\sigma^*) = 0 \forall j
\end{aligned}$$

TODO. Will complete from Solan...

Thus, by Brower's, we are done.

2.6. Two Player Non-Zero Sum Games

Definition

Given players 1, 2 and strategy sets S_1, S_2 and payoff matrices R for player 1 and C for player 2. A mixed strategy (x^*, y^*) is Nash Equilibrium if and only if

$$\begin{aligned}
x^{*T} R y^* &\geq (R y^*)_i \forall i \in [n] \\
x^{*T} C y^* &\geq (x^{*T} R) \forall i \in [n]
\end{aligned}$$

where the LHS are expected payoffs. Also (M_i) is the i^{th} position in the matrix. We could also write this as $e_i^T R y^*$ and $x^{*T} R e_i$ respectively.

Notice, if one switches say R for $R' = R + a$ that is

$$\begin{bmatrix} r_{1,1} + a & \dots & r_{1,m} + a \\ r_{2,1} + a & \dots & r_{2,m} + a \\ \vdots & \vdots & \vdots \\ r_{n,1} + a & \dots & r_{n,m} + a \end{bmatrix}$$

Claim 2.6. Set of NE for (R', C) are same as that of (R, C) . This implies, by symmetry, that $(R, C + b)$ has same NEs as (R, C) .

Claim 2.7. Set of NE remain unchanged for $(\frac{1}{a}R, C)$ wrt (R, C) and similarly the set of NE remains unchanged for $(R, \frac{1}{b}C)$ wrt (R, C) .

This allows us to only deal with $R_{i,j}, C_{i,j} \in [0, 1]$.

2.6.1. NE by Support enumeration

Let $S \subseteq [n]$, $T \subseteq [m]$ be set of strategies in the support of a NE that is $x_i > 0 \quad \forall i \in S$, $y_i > 0 \quad \forall i \in T$.

We will write a LP for this, named **LP[S,T]**.

$$\begin{aligned} \text{probability constraints} & \begin{cases} x_i \geq 0 & \forall i \in [n] \\ y_i \geq 0 & \forall i \in [n] \\ \sum_{i=1}^n x_i = 1 \\ \sum_{i=1}^m y_i = 1 \end{cases} \\ \text{NE Constraints} & \begin{cases} (Ry)_i \geq (Ry)_j & \forall i \in S, j \in [n] \\ (x^T C)_i \geq (x^T C)_j & \forall i \in T, j \in [m] \end{cases} \end{aligned}$$

Claim 2.8. A solution to the LP say (x^*, y^*) is a NE.

Proof. If x^*, y^* is not a NE then one of the players has an incentive to deviate to another pure strategy, which violates the NE constraints ■

NE ALGORITHM BY SUPPORT ENUMERATION

```

1  for each  $S \subseteq [n], T \subseteq [m]$ :
2    |   if  $\mathbf{LP}[S, T]$  has a feasible solution  $(x, y)$ 
3    |   |   return  $(x, y)$ 

```

This clearly has a $2^{n+m} \text{poly}(n+m)$ running time.

This is clearly not polytime. As we shall see, that the solution lies in complexity class PPAD and is hence, believed not to be possible in polytime.

2.6.2. Lemke-Howson Algorithm

Characterization of NE for this algorithm is using the translation and scaling transforms and using the fact that the max achievable payoff is always 1 and the payoffs are ≥ 0 .

We say $\left(\frac{x^*}{\sum^n x_i^*}, \frac{y^*}{\sum^m y_i^*}\right)$ is a NE if:

$$\begin{aligned} x_i^* = 0 & \quad \text{or} \quad (Ry^*)_i = 1 & \forall i \in [n] \\ y_i^* = 0 & \quad \text{or} \quad (x^{*T} C)_i = 1 & \forall i \in [m] \end{aligned}$$

The idea is that instead of the payoff's being variables, we can sort of normalize them to 1 and then make a system of linear equations using it. This will

Definition: Polytope, Polyhedron, Half-Space, Vertex

A bounded **polyhedron** is **polytope**.

Intersection of **half-spaces** is called a **polyhedron**.

Given a n dimensional plane, the plane can be split into two parts by a $n-1$ dimensional hyper-plane. This is called **half-space**. For example $ax + b > 0$ defines a half-space in 2D while $a_1x_1 + a_2x_2 + a_3x_3 \geq 0$ defines a half-space in 3D.

Vertex/Extreme Points in \mathbb{R}^k is the intersections of k linearly independent hyper-planes.

Consider the polytope:

$$P = \left\{ (x, y) \mid x_i \geq 0, y_i \geq 0, (Ry)_i \leq 1, (x^T C)_i \leq 1 \right\}$$

We are assuming non-degeneracy here.

Definition: Non-Degeneracy

Any set of $> m + n$ constraints do not meet at one point.

While the polytope itself doesn't define nash equilibria but some of it's vertex do. We will hop from vertex to vertex upto some condition to get nash equilibria.

At Nash equilibrium, by the stability condition, n of the equations in

$$\left\{ x_i \geq 0, (Ry)_i \leq 1 \right\}$$

and m of the equations in

$$\left\{ y_i \geq 0, (x^T C)_i \leq 1 \right\}.$$

Thus, $m + n$ many equalities the NE must have. That implies NE is a vertex of P .

The vertex $x, y = (\vec{0}_n, \vec{0}_m)$ is not a NE. This is called an artificial equilibrium.

A vertex $x > 0, Ry_i < 1$ but $x_2 = 0, R_2 y = 1, \dots$ is not a NE. Strategy 1 has +ve prov but not max payoff.

Let (\hat{x}, \hat{y}) be a vertex. Define a set of labels for each vertex. Define a set of labels for each vertex (\hat{x}, \hat{y}) has a label $i \in [n]$ if either $x_i = 0$ or $R_i y = 1$. Also (\hat{x}, \hat{y}) has a label $n + j$ if $y_j = 0$ or $x^T C^{(j)} = 1$.

Claim 2.9. (\tilde{x}, \tilde{y}) is a NE if and only if H has all $m + n$ labels and $(\tilde{x}, \tilde{y}) \neq (\vec{0}_n, \vec{0}_m)$
Under the non-degeneracy assumption.

Goal To find $(\tilde{x}, \tilde{y}) \neq (\vec{0}_n, \vec{0}_m)$ and has all the labels.

So let's start at $(x_0, y_0) = (\vec{0}_n, \vec{0}_m)$.

Fix label J , relax the constraint $x_1 = 0; x_2 = \dots = x_n = 0, y_1 = \dots = y_m = 0, x_1 > 0$.

This new vertex, say (x_1, y_1) doesn't have label 1, but has a duplicate label $\in \{n + 1, \dots, n + m\}$, say $k = 1$.

We relax the corresponding to the old duplicate label and proceed.

At any point, say t if (x_t, y_t) have all the labels, output it as a NE.

Claim 2.10. For any $i \neq j$, $(x_i, y_i) \neq (x_j, y_j)$

Proof. (x_t, y_t) has all labels and (x_0, y_0) is $(0, 0)$.

All intermediate vertices (x_i, y_i) with duplicate labels k_i have only two neighbors, corresponding to the duplicates. With only two ways of coming, there are only two ways of going, ensuring that we can never repeat vertices. ■

Corollary 2.11. *Every game has odd number of nash equilibrium.*

Corollary 2.12. *The odds in a mixed nash equilibrium are rational.*

2.7. Mixed Nash Equilibrium in 2 Player Games

$N = \{1, 2\}, S_1 = [n], S_2 = [m]$ with payoff matrices R, C .

Characterization of MNE is

P : set of points $(x, y) \in \mathbb{R}^{m+n}$ s.t.

$$x \geq 0 \forall i \in [n]$$

$$R_i y \leq 1 \forall i \in [n]$$

$$y_j \geq 0 \forall j \in [m]$$

$$x^T C_j \geq 0 \forall j \in [m]$$

$(x, y) \in P$ is a MNE if and only if

$\forall i \in [n]$, either $x_i = 0$ or $Ry_i = 1$: Label i cases

$\forall j \in [m]$, either $x_j = 0$ or $x^T C_j = 1$: Label $n + j$ cases

(x, y) has all labels in $[m + n]$.

Non-Degeneracy assumption \Rightarrow exactly $m + n$ labels for any vertex of P .

Note $(0, 0)$ has all labels but is not a NE.

LEMKE-HAWSON ALGORITHM

- 1 Start at $(0, 0)$, fix label 1.
 - 2 Relax the constraint \equiv label 1 (ie $x = 0$)
 - 3 Goto next vertex $(x^{(1)}, y^{(1)})$
 - 4 └ Either all labels exist \Rightarrow MNE
 - 5 Or a duplicate label k exists
 - 6 └ Relax the “old” constraint \equiv label k
 - 7 Goto $(x^{(2)}, y^{(2)})$
 - 8 so on...
 - 9 Let $(x^{(0)}, y^{(0)}), \dots, (x^{(t)}, y^{(t)})$ be the visited vertices.
-

LEMKE-HAWSON ALGORITHM

10 All $(x^{(i)}, y^{(i)}), 0 < i < t$ miss label 1.

This algorithm is clearly exponential time as our polytope may have exponential vertices.

The edges we visit form a graph, specifically a path where the first and last vertex have degree 1 and the other vertices have degree 2.

2.8. Complexity of algorithm

As with most problems in computer science, we want to know how hard it is. After all after being saddened by not getting a polytime algorithm, we want to show it is unlikely that anyone else will³.

The problem in showing that it is NP lies in the definition of NP itself.

Definition: NP

NP is a class of decision problems with yes/no answer and there exists a polytime verifiable certificate for a yes answer.

Our problem is not a decision problem. So what do we do?

Definition: Functional NP (FNP)

If there is a certificate (solution), output one.

Note, the certificate should be polytime verifiable.

Clearly, MNE \in FNP. So can we show MNE is FNP-Complete.

Theorem 2.13. *If MNE is FNP-Complete, $NP = co-NP$*

Side Note

It is believed that $NP \neq co-NP$, not as strongly as $P \neq NP$ but strongly enough. It is still open nonetheless.

Example

Take your favorite NP-complete problem. The instructor took Hamiltonian Path. Let's there be a reduction from Hamiltonian Path to MNE.

³Unless $P = NP$ which is still unlikely.

Input will be a graph G in which we want to compute the Hamiltonian path.

Reduction will be something that takes a graph G and converts it into an instance of 2-Player Game Γ with the property:

G has a Hamiltonian path $\iff \Gamma$ has a MNE which maps back to “yes”

G has no Hamiltonian path $\iff \Gamma$ has a MNE which maps back to “no”

Let's take a slightly less ambitious goal.

Definition: TFNP

Total FNP is the class of FNP problems that always have a solution.

This still doesn't help us out as there are no known TFNP problems.

So we go down to PPAD, a class contained in TFNP.

Definition: PPAD

Polynomial Parity Argument Directed version is a complexity class defined by a canonical problem called the **end of line** problem.

Definition: End of Line

Given G possibly exponential sized graph with polytime algorithm (or circuit) to determine neighbours of a given vertex.

Every vertex has in-degree ≤ 1 and out-degree ≤ 1 .

Q: Given a source, find a sink (any sink, not the one corresponding to the source!)

Theorem 2.14. NME is PPAD complete

High Level Proof. NME is in PPAD by the Lemke-Hawson algorithm. We do this reduction in 2 steps.

1. Reduce End of Line to (approx) Brouwer's Fixed Point.
2. Reduce approx Brouwer Fixed Point to approx MNE.

We will do the first part by an intermediate step called **Sperner Lemma**.

Sperner's Lemma 2.15. *Given a lattice as an input with the lattice points colored in three colors with every boundary being forbidden to use one color. The intermediate vertices can have any color (from the 3).*

We define a triangle as 3 points in the same cell. That is a $1, 1, \sqrt{2}$ right angle triangle where lengths are unit.

Sperner's Lemma states that there exists a triangle with all vertices having distinct colors.

Definition: Sperner's Lemma Problem

Given a lattice as an input with the lattice points colored in three colors with every boundary being forbidden to use one color. The intermediate vertices can have any color (from the 3).

We define a triangle as 3 points in the same cell. That is a $1, 1, \sqrt{2}$ right angle triangle where lengths are unit.

Find: a triangle with all vertices having distinct colors.

We can solve Brouwer Fixed Point on $f : [0, 1]^2 \rightarrow [0, 1]^2$, we want x s.t. $f(x) = x$.

We can just declare $0^\circ - 120^\circ$ as red, $120^\circ - 240^\circ$ as yellow and $240^\circ - 360^\circ$ as blue and color everything by this. We will use Sperner to find the approximate fixed point.

So our flow chart is

End of Line \rightarrow Sperner Problem \rightarrow Approx Brouwer Fixed Point
 \rightarrow Arithmetic Circuit \rightarrow Game \rightarrow MNE

■

2.9. Alternatives to Nash Equilibrium

Some issues with Nash Equilibria are:

- It is hard to compute (unless $\text{PPAD} \subseteq \text{P}$, which is unlikely)
- Many Nash Equilibria may exist, it is hence difficult to predict players behavior.
- Payoff at Nash Equilibrium may be much smaller than optimum (cost of anarchy)

Remark

For example, in Prisoner's Dilemma:

		P2	
		<i>C</i>	<i>D</i>
P1	<i>C</i>	−5, −5	−1, −10
	<i>D</i>	−10, −1	−2, −2

The optimal cases is $-2, -2$ but the Nash equilibrium is $-5, -5$ which makes the cost of anarchy 2.5 and is not desirable from players point of view.

Problem 2.16. Can we circumvent the hardness of Nash Equilibrium by computing say the ϵ -Nash Equilibrium?

Definition: ε -Nash Equilibrium

$(\sigma_1^*, \dots, \sigma_n^*)$ is a ε -NE if

$$\forall i \in [n] \quad \mu_i(\sigma_i^*, \sigma_{-i}^*) \geq \mu_i(\sigma_i', \sigma_{-i}^*) - \varepsilon \quad \forall \sigma_i' \in \Delta_i$$

Unfortunately, this is also PPAD complete for all ε as our proof for PPAD-hardness for NE uses approximate Brouwer's fixed point and more approximation won't really allow for any ease in computation.

Theorem 2.17. *Around each NE, there are ε -NE but converse may not hold as \exists games where an ε -NE is 'far' from any NE.*

2.9.1. Correlated Equilibria**Example**

		P2	
		<i>C</i>	<i>M</i>
P1	<i>C</i>	2, 1	0, 0
	<i>M</i>	0, 0	1, 2

The NE are clearly *C* and *M* pure and $((\frac{2}{3}, \frac{1}{3}), (\frac{1}{3}, \frac{2}{3}))$ mixed. The expected payoffs of the equilibrium are:

$$(2, 1), (1, 2), \left(\frac{2}{3}, \frac{2}{3}\right)$$

respectively where the mixed equilibrium is worse for both as they are assigning some probability to the undesirable *CM*, *MC* options.

These are called Correlated equilibria.

Definition: Correlated Equilibrium

Algorithm 2.18. . Given n -players $[n]$, strategy sets S_1, \dots, S_n .

- A coordinator declares a probability distribution on $S_1 \times \dots \times S_n$.
- The coordinator 'privately' samples the joint distribution for a strategy combination.
- The coordinator tells player i to play s_i

Note: The player may choose to or not to heed the advice.

A distribution D on $S_1 \times \dots \times S_n$ is a CE if $\forall i, \forall b_i, b_i' \in S_i$

$$\mathbb{E}_{s \sim D}[\mu_i(b_i, s_{-i}) \mid s_i = b_i] \geq \mathbb{E}_{s \sim D}[\mu_i(b'_i, s_{-i}) \mid s_i = b_i]$$

Basically, the player is better off heeding the advice of the coordinator given everyone else is heeding the advice of the coordinator.

Example : Chicken

		P2	
		<i>D</i>	<i>C</i>
P1	<i>D</i>	0, 0	5, 1
	<i>C</i>	1, 5	4, 4

This game comes from a dumb things teenagers did ‘back in the day’⁴ where they would drive towards each other at some speed and the first person to hit the breaks (Chicken out) would lose. The issue is if two hot headed (or stupid, albeit it is hard to tell the difference) drive into each other.

We can see the Nash Equilibria are $(A, B), (B, A), (\frac{1}{2}A, \frac{1}{2}B)$ whereas we can have a CE of $\frac{1}{3}(A, B), \frac{1}{3}(B, A), \frac{1}{3}(B, B)$.

As John Green once said: “Turtles all the way down!”, we will say Equilibria all the way down.

Dominant Equilibria \subseteq Weakly Dominant Equilibria \subseteq Nash Equilibria \subseteq Corelated Equilibria

where the equilibria from Nash onwards are gurenteed to exist and Corelated is easy to compute.

Theorem 2.19. *CE form a convex set*

Algorithm 2.20. We can compute CE using LP:

$$\begin{aligned} \sum_j A_{ij} p_{ij} &\geq \sum_j A_{i'j} p_{ij} \\ \sum_j B_{ij} p_{ij} &\geq \sum_i B_{ij'} p_{ih} \\ \sum_{i,j} p_{ij} &= 1 \\ p_{ij} &> 0 \end{aligned}$$

⁴Nowadays, teenagers have better (read safer) things to do in their time. Unfortunately, social media fame, friendship graphs, min-maxing for university/NEET-JEE/Olympiads etc are much harder to study.

3. Mechanism Design

Definition: Mechanism Design

Till now we have been given players, game and strategies and we want to find outcome via the equilibrium.

We will now do mechanism design where we are given the players and our ideal outcome and we want to come up with a game where the equilibrium is in line with our outcome.

Example

In making a cricket tournament, we don't want players to lose on purpose to get an easier opponent later and have better medal odds.

A failure of this was the 2012 London Olympics where the women's badminton doubles was structured so badly that both teams were trying to lose in an extremely hard to watch match. Search "Wang Xiaoli / Yu Yang (CHN) vs Ha Jung-eun / Kim Min-jung (KOR)" if you want to see the match, the referee being tense and the crowd booing.

They were later disqualified for unsportsmanlike behavior. Although, trying to win by all legal means is sportsmanlike in my book. Maybe they should have set the tournament better to never let this be incentivized in the first place.

Mechanism design is used to design tournaments, voting schemes and schemes to divide stuff.

4. Fair Division

Definition: Division

$\langle X_1, X_2, \dots, X_n \rangle$ a division of $[0, 1]$ if

$$X_i \subset [0, 1],$$

$$X_i \cap X_j = \emptyset$$

$$\bigcup_{i=1}^n X_i = [0, 1]$$

Definition: Envy Free

Given n agents and their valuation functions v_i , and we want to divide a cake $[0, 1]$ then $\langle X_1, X_2, \dots, X_n \rangle$ a EF division if

$$\forall i, j \in [n]; v_i(X_i) \geq v_i(X_j)$$


Problem 4.1. Envy Free Cake Division on 2 agents**Solution**

This is rather easy. We ask one agent to cut and other to choose.

Another idea is that I keep moving the knife and when an agent is satisfied say at a , I cut the cake there and give them the piece $[0, a]$ as

Problem 4.2. Envy Free Cake Division with 3 agents

Problem 4.3. Given an apartment with n rooms and rent k , we want to divide the apartment between n roommates with valuation functions $v_i : [n] \times [0, k] \rightarrow [0, 1]$.

4.1. Detour: Sperner's Lemma**Definition: Sperner Coloring on Triangles**

Given a triangle ABC divided into triangles, a coloring of vertices such that

- Each of the three vertices A , B , and C of the initial triangle has a distinct color.
- The vertices that lie along any edge of triangle ABC have only two colors, the two colors at the endpoints of the edge. For example, each vertex on AC must have the same color as A or C .

Sperner's Lemma in Triangles 4.4. *Given any Sperner colored triangulation, there exists odd number of fully colored elementary triangle (has all three colors.)*

Corollary 4.5. *Given any Sperner colored triangulation, there exists atleast one of fully colored elementary triangle.*

Definition: Simplex

A 0 simplex is just a point.

A 1 simplex is a line.

A 2 simplex is a triangle.

A 3 simplex is a tetrahedron.

So on and so fourth.

Definition: Facet

Choosing n vertices from a $n + 1$ -simplex is called a facet.




Figure 1: Sperner Tetrahedron

Definition: Sperner Coloring

Given a n simplex divided into n simplices, we call a coloring Sperner if:

- Each of the boundary vertices of the initial simplex has a distinct color.
- The vertices that lie along any facet i (named after the boundary vertex not in the facet) doesn't have color i .

Sperner's Lemma 4.6. *Any Facet of a sperner colored simplex is also sperner colored, also, number of rainbow elementary n -simplices is odd.*

Proof.

Definition: Door

Door is a facet of elementary simplex if it has labels $\langle 1 \dots n \rangle$. For a 3 simplex (triangle), it is edges with vertices colored $(1, 2)$.

Lemma 4.7. *We have an odd number of doors per facet.*

We will now induct on n .



Figure 2: Proof for 1D case

(B) For 1 simplex, We are obviously done as 1 – 2 rainbow will be there.

(S) Let Sperner hold for $n - 1$ simplex.




Figure 3: Visualization for the inductive step in 2D

For an n simplex, The n facet is $n - 1$ simplex and by induction, we will have odd number of doors via Sperner.

We can now enter via doors and lock them behind us. If we end up in a room, it is rainbow.

If we exit the building, we will close the door we left through and close it behind us.

If some rainbow room is unreachable, we claim that if we were to start there, we would end up walking to another unreachable rainbow room. Thus, the unreachable rooms come in pairs.

Thus, we have odd number of rainbow rooms, proving the Sperner's Lemma. ■

4.2. Back to Fair Division

4.2.1. Cake-Cutting Problem

We will solve a general version of Problem 4.2 also called the cake-cutting problem. For n agents, We consider any $n - 1$ vertical cuts parallel to the side of the cake, dividing the cake into n pieces. Such a way of cutting is called a cut-set.

Idea

Observe that in a particular cut-set, if we ask all players which part do they prefer and get all different answers, then that cut-set gives an envy-free division (because every one gets the piece he/she prefer the most). We will use Sperner's lemma to prove that there must exist such cut-set, as well as finding one.


The idea is that Sperner allows us to guarantee that there is a triangle part with all vertices of different color. So given continuity aka small change in cut-set causes small change in preference; we can make a Sperner coloring of the linear programming simplex and get an approximate. That's exactly what the solution is.

Solution. Consider every possible cut-set that divides the cake into n (maybe empty) pieces. Let x_i be the proportion of i th leftmost part. We have $x_1 + x_2 + \dots + x_n = 1$ and $x_i \geq 0$ for every $i = 1, 2, \dots, n$.

In the n -dimensional space, consider a polytope formed by linear programming $x_1 + x_2 + \dots + x_n = 1$ and $x_i \geq 0$ for every $i = 1, 2, \dots, n$. The resulted polytope is a regular $(n - 1)$ -simplex.

Then, we triangulate that simplex into smaller regular $(n - 1)$ -simplices with each having side length less than ε , for a small enough ε , as well as writing numbers $1, 2, \dots, n$ on the vertices in a way that every elementary $(n - 1)$ -simplex has vertices with all different numbers.

Example



In that figure, the resulted polytope is a regular triangle (regular 2-simplex). We then divide it into k^2 smaller regular triangles, for some integer k as big as we want. Then, we write a number on each vertex $1, 2, 3$ in a cyclic order such that every elementary triangle has all three vertices with different numbers.

Then, at each vertex with coordinates (x_1, x_2, \dots, x_n) with number i , we ask player i that which piece of cake that he/she prefers if the sizes of pieces of cake is x_1, x_2, \dots, x_n , respectively. We then color that vertex according to the answer.

Example

We ask player 1 that if the sizes of three pieces are 0, 0, and 1, respectively, which piece does he/she prefer. We then color that vertex according to the answer.

Observe that, in each of the k -dimensional face of the polytope, one or more of the coordinates must be zero.

Example

The second coordinate along the left edge of the triangle is always zero, hence, no one chooses that piece.

As, no people prefer the piece with size zero, so the color of vertices on each face must be the same as one of the corners of that face. Therefore, the color labeling of vertices in the simplex is a Sperner labeling.

By Sperner's lemma, there must be at least one elementary $(n - 1)$ simplex that has vertices with different colors. We then divide the cake by the cut-set represented by any interior point of that $(n - 1)$ -simplex.

Since the size of each elementary simplex is less than ε , by the continuous preference assumption, all people will be satisfied with that cut-set. ■

The algorithm this solution leads to is named **Simmons' Algorithm** as it was developed by Forrest Simmons in 1980.

4.2.2. Harmonious Rent Problem


We will now look into solving Problem 4.3 using Sperner's lemma.

One possible attempt is to consider every possible assignment of price to each bedroom. i.e. set the price of room i to be x_i such that $x_1 + x_2 + \dots + x_n = S$, when S is the total price of the apartment, and $x_i \geq 0$ for every $i = 1, 2, \dots, n$.

Then, we consider the $(n - 1)$ -simplex obtained from the linear programming and color it like in the cake-cutting problem.

However, this problem is different from the cakecutting problem in one aspect. In the cake-cutting problem, the bigger the piece of cake is, the higher chance people will want

it; however, in this problem, the higher the price of a room is, the lower chance people will want that room. This would mean our simplex could look like:



Idea

If we could somehow turn every corner into a face and every face to a corner, we would be done.

Can that be done?

Yes. We transform each k face to $n - k - 1$ face (for a n simplex).




Figure 4: We have marked the vertices, points on line and colored an elementary cell to show how it transforms.

Let the triangulation of S^* inherit a labelling via this correspondence with S . One may now verify that the labelling of S^* is a Sperner labelling! Hence, there exists a fully labelled elementary simplex of S^* , which corresponds to a fully labelled elementary simplex of S , as desired. We can now present the solution.

Solution

A constructive algorithm is obtained by following “trap-doors” in Sperner’s lemma. Choose an ε smaller than the rental difference for which housemates wouldn’t care (a penny?). Following trap-doors corresponds to suggesting pricing schemes and then asking various players, “Which piece would you choose if the rooms were priced like this?” Once a fully labelled elementary simplex is found, any point inside it corresponds to an ε -approximate rent-partitioning.

However, we can also solve this without ever needing to go through this dualization process.

Solution

Each person first pays the full rent price and then assigns to each room a “rebate” price that a person who takes that room will get. Then, we can use Simmons’ Algorithm to assign rooms to people.

4.3. Fair Division of indivisible goods

TODO. *Lecture on 19 th*

As we have shown, in divisible case, EF allocation always exists.

In indivisible goods, EF might not always exist. For example, **2 agents 1 good**.

However, we can do EF1 in P.

Definition: EF1

An allocation is EF1 if for any agents $i, j \in [n]$

$$\exists g \in A_j \text{ s.t. } v_i(A_i) \geq v_j(A_i - g)$$

that is agents don’t envy each other upto removal of one good (or addition of one chore.)

Definition: Pareto

An allocation $A \succ B$ if

$$\begin{aligned}\forall i \in [n] \quad v_i(A_i) &\geq v_i(B_i) \\ \exists j \in [n] \quad v_j(A_j) &\not\geq v_j(B_j)\end{aligned}$$

Definition: Pareto Optimal (PO)

A is Pareto optimal if there is no B such that $B \succ A$.

Basically, A is PO if you cannot make some agents better off without making anyone worse off.

Note, PO always exists and the proof is by the fact that giving everything to the same agent does achieve that.

Definition: Social Welfare

Social welfare of allocation A is

$$SW(A) = \sum_{i \in [n]} v_i(A_i)$$

Theorem 4.8. *If*

$$A \in \arg \max_{B \in \Pi_m(m)} SW(B)$$

then A is PO.

Proof. If A is not PO, we can just switch the goods around and end up with higher valuation for some agents and hence, a higher social welfare which contradicts the maximality of social welfare. ■

So can we get a **EF1** + **PO**? Trying to do pareto switches to EF1 doesn't work as:

$$\begin{pmatrix} A : 1 & 1 & \dots & 1 \\ B : 1 & 1 & \dots & 1 \\ C : 0 & 0 & \dots & 0 \end{pmatrix}$$

here the allocation $A_1 = \{1, 2, \dots, 10\}$, $A_2 = \{11, 12, \dots, 20\}$, $A_3 = \{21, 22, \dots, 30\}$ is EF but making pareto switches will make it extremely envy prone.

Does the round robin algorithm work?

$$\begin{pmatrix} A : \boxed{2} & 2 & \boxed{2} & 2 \\ B : 10 & \boxed{1} & 1 & \boxed{1} \end{pmatrix}$$

This is not pareto optimal as switching $10 \leftrightarrow 2$ and $2 \leftrightarrow 1$ can be done.

Does Envy cycle agent elimination work?

$$\begin{pmatrix} A : \boxed{10} & 5 & \boxed{1} & 1 \\ B : 1 & \boxed{1} & 5 & \boxed{10} \end{pmatrix}$$

but this is not PO as switching $5 \leftrightarrow 1$ and $1 \leftrightarrow 5$ can be done.

So does a EF1 + PO allocation even exist?

Definition: Nash Social Welfare

$$\text{NSW}(A) = \left(\prod_{i \in [n]} v_i(A_i) \right)^{\frac{1}{n}}$$

Definition: Nash Optimal

We say A is Nash Optimal if it is the largest set of agents S who can simultaneously get positive values and take such a division which maximizes the geometric mean among S .

Working in an instance where $m \geq n$ where $\forall i \in [n], \forall g, v_i(g) > 0$.

Lemma 4.9. *If A is Nash optimal then:*

1. A is PO
2. A is EF1

Proof of (1).. Suppose not. Say $B \succ A$ then the set of people with non-zero utility in B , say $S' \supseteq S$ which is the set of people with non zero utility in A . Thus,

$$\text{NSW}(B) = \prod_{i \in S'} v_i(B_i) > \prod_{i \in S} v_i(B_i) \geq \prod_{i \in S} v_i(A_i)$$

which is a contradiction. ■

Proof of (2).. Suppose A is Nash optimal but not EF1.

That is $\exists i, j$ such that

$$\forall g \in A_j v_i(A_i) < v_i(A_k - g)$$

Let

$$g^* = \arg \min_{\substack{g \in A_k \\ v_i(g) > 0}} \frac{v_k(g)}{v_i(g)}$$

Construct a new allocation with everything same but $v_i(B_i) = v_i(A_i) + v_i(g^*)$ and $v_j(B_j) = v_j(A_j) - v_j(g^*)$.

We claim

$$\begin{aligned}
& \frac{\text{NSW}(B)}{\text{NSW}(A)} > 1 \\
& \iff v_i(B_i)v_j(B_j) > v_i(A_i)v_j(A_j) \\
& \iff [v_i(A_i) + v_i(g^*)][v_j(A_j) - v_j(g^*)] > v_i(A_i)v_j(A_j) \\
& \iff v_i(A_i)v_j(A_j) + v_i(g^*)v_j(A_j) - v_i(A_i)v_j(g^*) - v_i(g^*)v_j(g^*) > v_i(A_i)v_j(A_j) \\
& \iff 1 + \frac{v_i(g^*)}{v_i(A_i)} - \frac{v_j(g^*)}{v_j(A_j)} - \frac{v_i(g^*)v_j(g^*)}{v_i(A_i)v_j(A_j)} > 1 \\
& \iff v_j(A_j) > \frac{v_j(g^*)}{v_i(g^*)}[v_i(A_i) + v_i(g^*)] \quad (\text{by Vishwa bhaiya claims this to be true})
\end{aligned}$$

From the choice of g^* ,

$$\begin{aligned}
& \frac{v_j(g^*)}{v_i(g^*)} = \theta \\
& \Rightarrow \frac{v_j(g)}{v_i(g)} \geq \theta \\
& \Rightarrow v_j(g) \geq \theta v_i(g)
\end{aligned}$$

Using this,

$$\frac{v_j(A_j)}{v_i(A_j)} \geq \theta = \frac{v_j(g^*)}{v_i(g^*)}$$

Using envy freeness,

$$\begin{aligned}
v_i(A_i) & < v_i(A_j) - v_i(g^*) \\
v_i(A_i) + v_i(g^*) & < v_i(A_j)
\end{aligned}$$

We can multiply these to get

$$v_j(A_j) > \frac{v_j(g^*)}{v_i(g^*)}[v_i(A_i) + v_i(g^*)]$$

leading to a contradiction and completing the proof. ■

While NSW is gurenteed to exist, it is NP hard to find and even hard to approximate. The question if EF1+PO can be done in poly time is open.